

3a. Logistic Regression - Serve Statistics

Using logistic regression, the three serve statistics are tested to see their correlation with the likelihood of winning a match on the ATP tour. Statistics for the winner and loser were created and a new variable was added (1 for a match won and 0 for a match lost). The two dataframes were then combined and are ready for the regression.

Table of Content

1. Load and Verify Data
2. Separate and Rename Columns
3. Combine the Dataframes
4. Train the Logistic Regression and Plot
5. Export Dataframe (for use on Tableau)

Step 1: Load and Verify Data

```
import pandas as pd
import os
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
import matplotlib.pyplot as plt
import seaborn as sns

#Set Path
path = r'/Users/tristansavella/Desktop/Important Things/Data
Analytics/CareerFoundry/Data Immersion/Achievement 6/Master Folder
ATP/02 Data'

#Import df_matchstats
df_matchstats = pd.read_pickle(os.path.join(path, 'Prepared
Data', 'df_matchstats.pkl'))

# Set the option to display all rows
pd.set_option('display.max_rows', None)
```

```
# Set the option to display all columns
pd.set_option('display.max_columns', None)
```

```
df_matchstats.head()
```

	tourney_id	Year	tourney_name	surface	tourney_level	
winner_id \						
119317	2000-301	2000	Auckland	Hard	A	103163
119318	2000-301	2000	Auckland	Hard	A	102607
119319	2000-301	2000	Auckland	Hard	A	103252
119320	2000-301	2000	Auckland	Hard	A	103507
119321	2000-301	2000	Auckland	Hard	A	102103

	winner_ioc	winner_name	winner_age	winner_rank
winner_ht \				
119317	GER	Tommy Haas	21.7	11.0
188.0				
119318	ESP	Juan Balcells	24.5	211.0
190.0				
119319	ESP	Alberto Martin	21.3	48.0
175.0				
119320	ESP	Juan Carlos Ferrero	19.9	45.0
183.0				
119321	USA	Michael Sell	27.3	167.0
180.0				

	loser_id	loser_ioc	loser_name	loser_rank
loser_ht \				
119317	101543	USA	Jeff Tarango	63.0
				180.0
119318	102644	ARG	Franco Squillari	49.0
				183.0
119319	102238	ESP	Alberto Berasategui	59.0
				173.0
119320	103819	SUI	Roger Federer	61.0
				185.0
119321	102765	FRA	Nicolas Escude	34.0
				185.0

	loser_age	best_of	round	minutes	w_#ServeGames	w_#aces
w_#dfs \						
119317	31.1	3	R32	108.0	17.0	18.0
4.0						
119318	24.3	3	R32	85.0	12.0	5.0
3.0						
119319	26.5	3	R32	56.0	8.0	0.0

0.0						
119320	18.4	3	R32	68.0	10.0	5.0
1.0						
119321	23.7	3	R32	115.0	13.0	1.0
2.0						

	w_#ServePoints	w_#1stServesIn	w_#2ndServePoints	w_#1stServesIn
119317	96.0	49.0	47.0	
51				
119318	76.0	52.0	24.0	
68				
119319	55.0	35.0	20.0	
63				
119320	53.0	28.0	25.0	
52				
119321	98.0	66.0	32.0	
67				

	w_#1stWon	w_%1stWon	w_#2ndWon	w_%2ndWon	w_bpSaved
119317	39.0	79	28.0	59	3.0
5.0					
119318	39.0	75	13.0	54	5.0
6.0					
119319	25.0	71	12.0	60	1.0
1.0					
119320	26.0	92	15.0	60	0.0
0.0					
119321	39.0	59	14.0	43	6.0
11.0					

	l_#ServeGames	l_#aces	l_#dfs	l_#ServePoints	l_#1stServesIn
119317	17.0	7.0	8.0	106.0	55.0
119318	12.0	5.0	10.0	74.0	32.0
119319	8.0	0.0	6.0	56.0	33.0
119320	10.0	11.0	2.0	70.0	43.0
119321	12.0	8.0	8.0	92.0	46.0

	l_#2ndServePoints	l_#1stServesIn	l_#1stWon	l_%1stWon
119317	51.0	51	39.0	70
29.0				
119318	42.0	43	25.0	78

18.0				
119319	23.0	58	20.0	60
7.0				
119320	27.0	61	29.0	67
14.0				
119321	46.0	50	34.0	73
18.0				

	l_%2ndWon	l_bpSaved	l_#bpFaced
119317	56	4.0	7.0
119318	42	3.0	6.0
119319	30	7.0	11.0
119320	51	6.0	8.0
119321	39	5.0	9.0

Drop columns with NaN values

```
df_matchstats.dropna(axis=1, inplace=True)
```

Verify initial columns

```
print("Initial Columns:", df_matchstats.columns.tolist())
```

```
Initial Columns: ['tournament_id', 'Year', 'tournament_name', 'surface',
'tournament_level', 'winner_id', 'winner_ioc', 'winner_name',
'winner_age', 'loser_id', 'loser_ioc', 'loser_name', 'best_of',
'round', 'w_#ServeGames', 'w_#aces', 'w_#dfs', 'w_#ServePoints',
'w_#1stServesIn', 'w_#2ndServePoints', 'w_#1stServesIn', 'w_#1stWon',
'w_#1stWon', 'w_#2ndWon', 'w_#2ndWon', 'w_bpSaved', 'w_#bpFaced',
'l_#ServeGames', 'l_#aces', 'l_#dfs', 'l_#ServePoints',
'l_#1stServesIn', 'l_#2ndServePoints', 'l_#1stServesIn', 'l_#1stWon',
'l_#1stWon', 'l_#2ndWon', 'l_#2ndWon', 'l_bpSaved', 'l_#bpFaced']
```

Step 2: Separate and Rename Columns

Create separate DataFrames for winner and loser statistics

```
winners = df_matchstats.copy()
```

```
losers = df_matchstats.copy()
```

Add a target variable to each DataFrame

```
winners['win'] = 1
```

```
losers['win'] = 0
```

Rename columns for losers to ensure uniqueness

```
losers.rename(columns={
    'l_#1stServesIn': '1stServesIn',
    'l_#1stWon': '1stWon',
    'l_#2ndWon': '2ndWon'
}, inplace=True)
```

```

# Rename columns for winners to ensure consistency
winners.rename(columns={
    'w_%1stServesIn': '1stServesIn',
    'w_%1stWon': '1stWon',
    'w_%2ndWon': '2ndWon'
}, inplace=True)

# Extract relevant columns
winners_renamed = winners[['1stServesIn', '1stWon', '2ndWon', 'win']]
losers_renamed = losers[['1stServesIn', '1stWon', '2ndWon', 'win']]

# Verify the column names after renaming
print("Winners Renamed Columns:", winners_renamed.columns.tolist())
print("Losers Renamed Columns:", losers_renamed.columns.tolist())

Winners Renamed Columns: ['1stServesIn', '1stWon', '2ndWon', 'win']
Losers Renamed Columns: ['1stServesIn', '1stWon', '2ndWon', 'win']

```

Step 3: Combine the DataFrames

```

# Combine the DataFrames
combined_df = pd.concat([winners_renamed, losers_renamed],
    ignore_index=True)

# Verify the combined DataFrame
print("Combined DataFrame Columns:", combined_df.columns.tolist())
print(combined_df.head())

Combined DataFrame Columns: ['1stServesIn', '1stWon', '2ndWon', 'win']

```

	1stServesIn	1stWon	2ndWon	win
0	51	79	59	1
1	68	75	54	1
2	63	71	60	1
3	52	92	60	1
4	67	59	43	1

Step 4: Train the Logistic Regression Model and Plot

```

# Define the features and target variable
features = combined_df[['1stServesIn', '1stWon', '2ndWon']]
target = combined_df['win']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(features, target,
    test_size=0.2, random_state=42)

```

```

# Train the logistic regression model
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)

LogisticRegression(max_iter=1000)

# Make predictions and evaluate the model
y_pred = log_reg.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(f'Classification Model Accuracy: {accuracy}')
print('Confusion Matrix:\n', conf_matrix)
print('Classification Report:\n', class_report)

Classification Model Accuracy: 0.7902606748760594
Confusion Matrix:
[[9794 2606]
 [2640 9972]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.79	0.79	0.79	12400
1	0.79	0.79	0.79	12612
accuracy			0.79	25012
macro avg	0.79	0.79	0.79	25012
weighted avg	0.79	0.79	0.79	25012

```

# Get the coefficients of the logistic regression model
coefficients = log_reg.coef_[0]
feature_names = features.columns

# Create a DataFrame for feature importance (coefficients)
importance_df = pd.DataFrame({'Serve Statistic': feature_names,
                              'Coefficient': coefficients})
importance_df = importance_df.sort_values(by='Coefficient',
ascending=False)
print(importance_df)


```

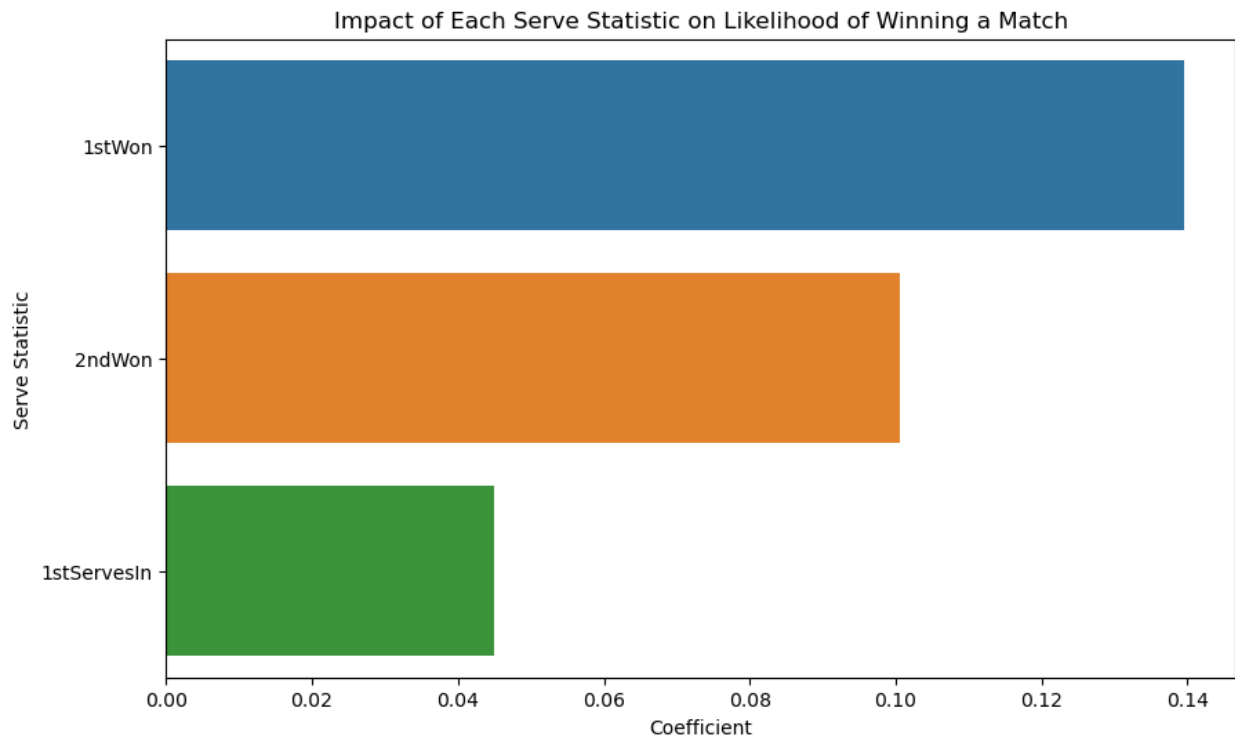
	Serve Statistic	Coefficient
1	1stWon	0.139635
2	2ndWon	0.100579
0	1stServesIn	0.044955

```

# Plot the feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x='Coefficient', y='Serve Statistic', data=importance_df)
plt.title('Impact of Each Serve Statistic on Likelihood of Winning a

```

```
Match')  
plt.show()
```



Step 5: Export Combined Dataset

```
combined_df.to_pickle(os.path.join(path, 'Prepared  
Data', 'combined_df.pkl'))  
  
#as CSV  
combined_df.to_csv(os.path.join(path, 'Prepared Data  
CSV', 'combined_df.csv'))  
  
importance_df.head()
```

	Serve Statistic	Coefficient
1	1stWon	0.139635
2	2ndWon	0.100579
0	1stServesIn	0.044955